

Linux Heap Introduction

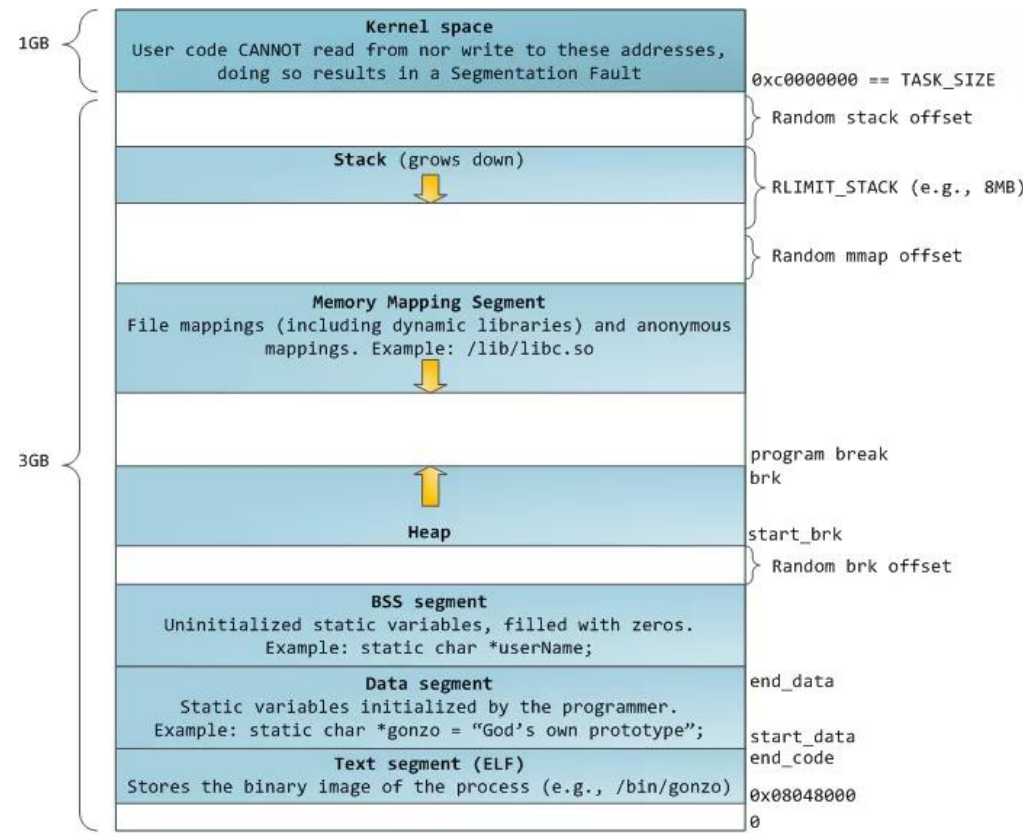
--w0lfzhang

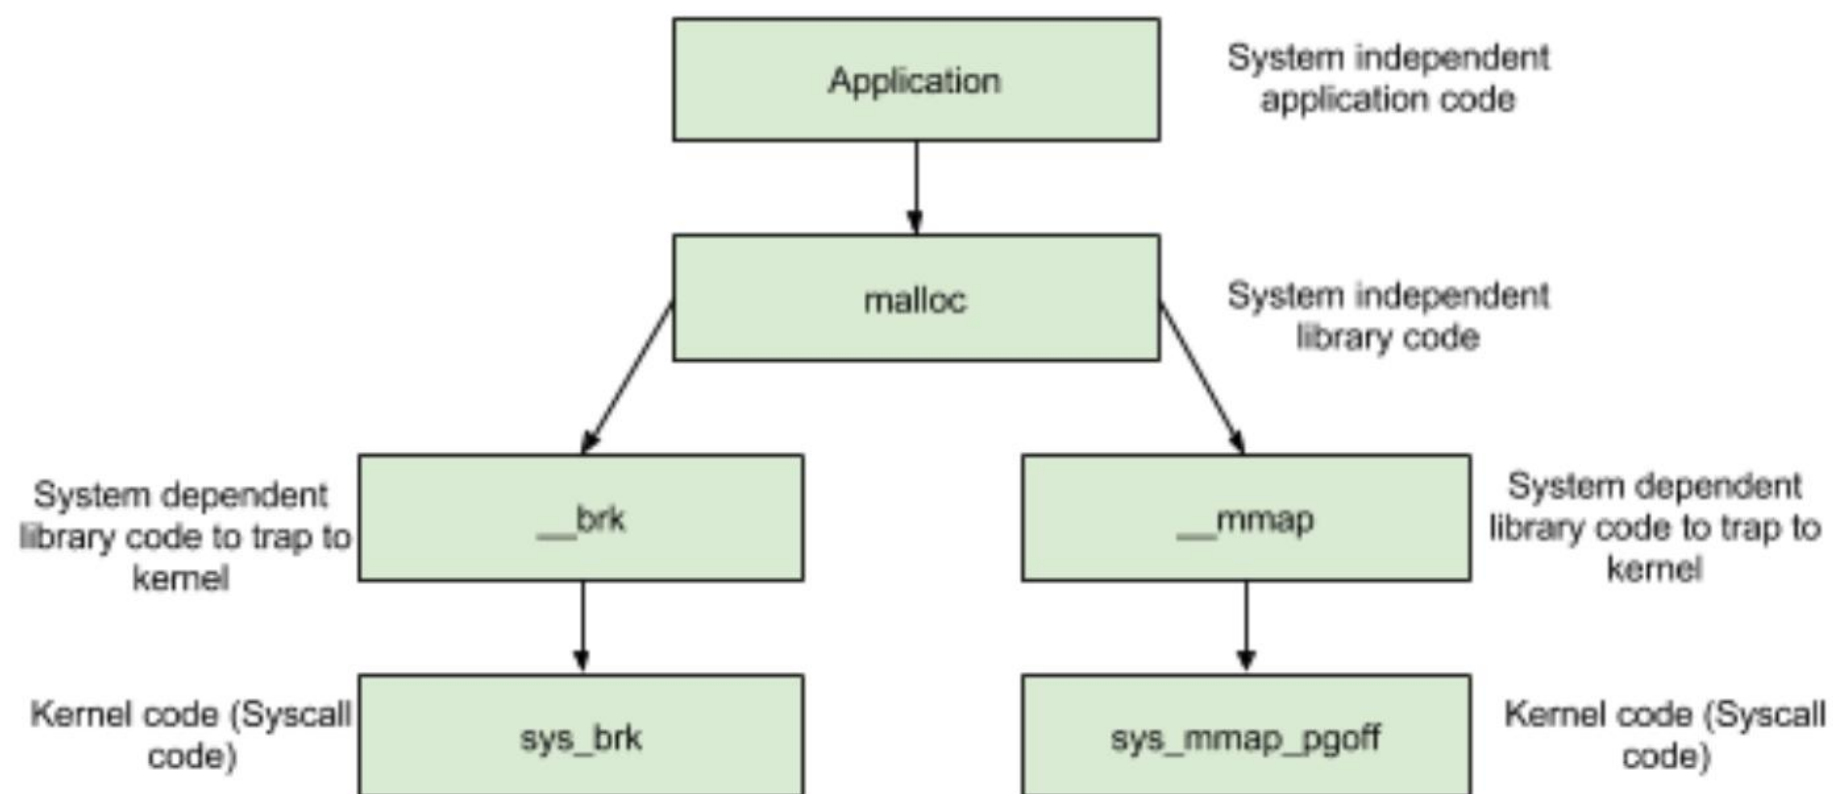
blog:w0lfzhang.me

- Linux Heap Structure
- Some Techs in CTFs

Linux glibc: ptmalloc2

- Linux's *malloc: obtaining memory by invoking either **brk** or **mmap** syscall.
- mmap: when needing large memory.(> the value system setted)





Some important data structures

- `heap_info`: heap header. arena can have multiple heaps(not including main thread)
- `malloc_state`: arena header. The process have just one `main_arena` and `non_main_arena`(maybe > 1). The thread just have one arena.
- `malloc_chunk`: chunk header. Fastbin, smallbin and largebin
- Main thread doesn't have the `heap_info` structure
- `main_arena` is a global variable

heap_info

- typedef struct _heap_info
- {
- mstate ar_ptr; /* Arena for this heap. */
- struct _heap_info *prev; /* Previous heap. */
- size_t size; /* Current size in bytes. */
- size_t mprotect_size; /* Size in bytes that has been mprotected
- PROT_READ|PROT_WRITE. */
- /* Make sure the following data is properly aligned, particularly
- that sizeof (heap_info) + 2 * SIZE_SZ is a multiple of
- MALLOC_ALIGNMENT. */
- char pad[-6 * SIZE_SZ & MALLOC_ALIGN_MASK];
- } heap_info;

malloc_state

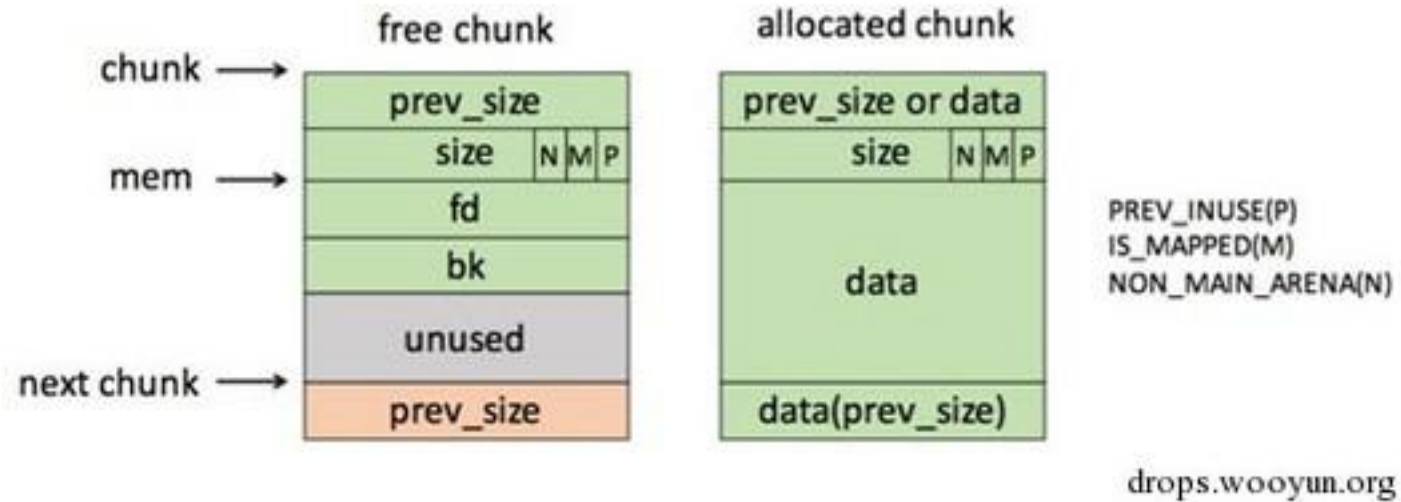
- struct malloc_state
- {
- /* Serialize access. */
- mutex_t mutex;
- /* Flags (formerly in max_fast). */
- int flags;
- /* Fastbins */
- mfastbinptr fastbinsY[NFASTBINS];
- /* Base of the topmost chunk -- not otherwise kept in a bin */
- mchunkptr top;
- /* The remainder from the most recent split of a small request */
- mchunkptr last_remainder;
-

- `/* Normal bins packed as described above */`
- `mchunkptr bins[NBINS * 2 - 2];`
- `/* Bitmap of bins */`
- `unsigned int binmap[BINMAPSIZE];`
- `/* Linked list */`
- `struct malloc_state *next;`
- `/* Linked list for free arenas. */`
- `struct malloc_state *next_free;`
- `/* Memory allocated from the system in this arena. */`
- `INTERNAL_SIZE_T system_mem;`
- `INTERNAL_SIZE_T max_system_mem;`
- `};`

malloc_chunk

- struct malloc_chunk {
 - INTERNAL_SIZE_T prev_size; /* Size of previous chunk (if free). */
 - INTERNAL_SIZE_T size; /* Size in bytes, including overhead. */
 - struct malloc_chunk* fd; /* double links -- used only if free. */
 - struct malloc_chunk* bk;
 - /* Only used for large blocks: pointer to next larger size. */
 - struct malloc_chunk* fd_nextsize; /* double links -- used only if free. */
 - struct malloc_chunk* bk_nextsize;
 - };

chunks

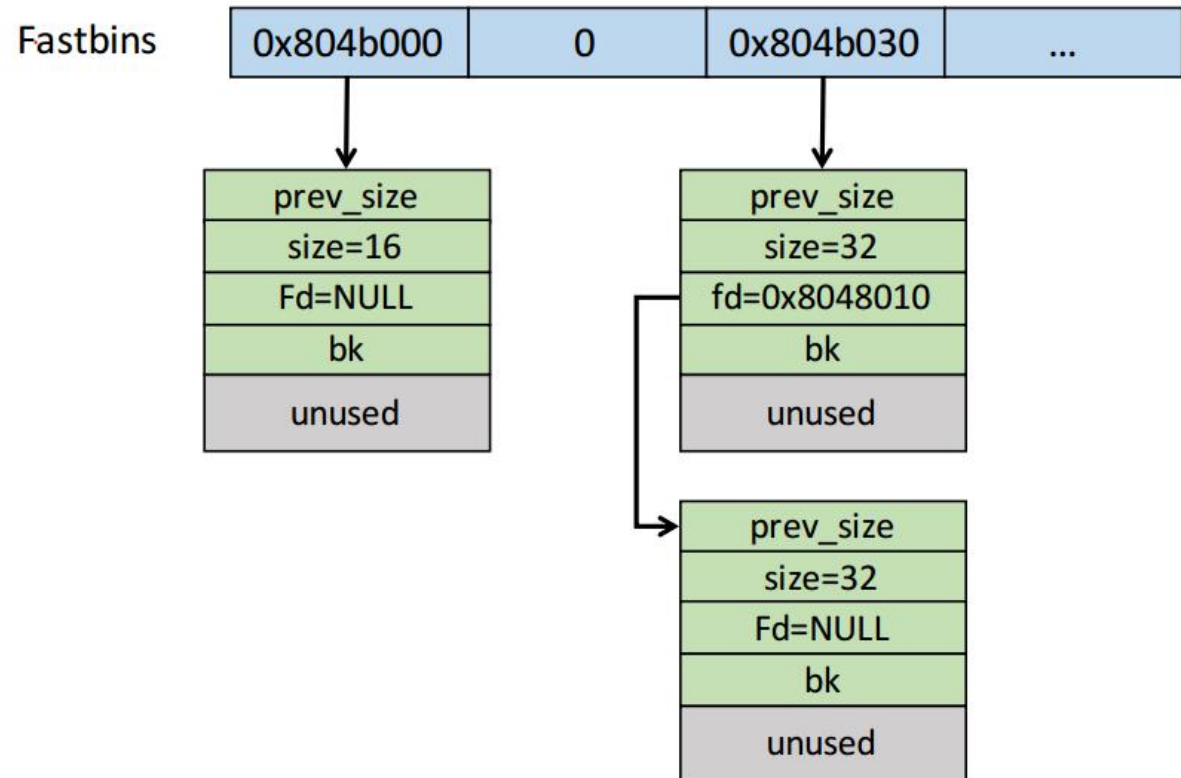


Bins

- `fastbinY[]`: just including fast bin
- `bins[]`: including unsorted bin, small bin and large bin
- bin 1: unsorted bin
- bin 2-bin 63: small bin
- bin 64-bin 126: large bin

Fast bin

- 16~64 bytes for x86_32
- 32~128 bytes for x86_64
- LIFO: Single Linked List
- No coalescing



Unsorted bin

- when the chunk was freed, the chunk was added in unsorted bin(not including fast bins)
- when mallocing a chunk, if the chunk's size(in unsorted bin) $>$ requested chunk's size, malloc from the chunk(by cutting it) and the rest becomes the last reminder chunk
- if the chunk's size $<$ requested chunk's size, malloc from top chunk and the chunk is linked into the corresponding bin

Small bin

- 16, 24, 32, ..., 508 bytes for x86_32
- FIFO: Circular double linked list
- Coalescing when next to each other
- After free(), check if the next chunk is freed; if so, coalesce them and add the new chunk into unsorted bin

Large bin

- ≥ 512 bytes(x86_32)
- the chunk's size can be **different** in the bin
- malloc: finding the **proper** chunk
- if =, mallocing directly
- if >, cutting it and malloc the requested size chunk and the rest is added into unsorted bin

Top chunk & Last Reminder chunk

- Top chunk: the top chunk of an arena
- Used to service user request when there is NO free blocks
- Extended using sbrk (main arena) or mmap (thread arena) syscall
- Last Reminder chunk: remainder from the most recent split of a small request

Some Useful Skills in CTFs

- Malloc Maleficarum:
 - house of force
 - house of spirit
 - house of lore
 - house of mind
 - house of prime
- blog: gbmaster.wordpress.com

unlink

- #define unlink(P, BK, FD)
- {
-
- BK = P->bk;
- FD = P->fd;
- FD->bk = BK;
- BK->fd = FD;
-
- }

- The check:
- `if (__builtin_expect (FD->bk != P || BK->fd != P, 0))`
- `malloc_printerr (check_action, "corrupted double-linked list", P);`


- **Solution:**

- Find a pointer ptr pointing to p(*ptr == p)
- `p->fd = ptr - 0x18(64bit)`
- `p->bk = ptr - 0x10`
- Finally, `p = ptr - 0x18`

double free

- free a pointer twice
- using unlink essentially
- p1 = malloc(n)
- p2 = malloc(m)
- free both of them
- a wild pointer p2

- then malloc(x)
- $x > n + 16(32\text{bit})$
- so as to make fake data to use the wild pointer p



- Understanding the heap by breaking it - Black Hat
- <http://www.blackhat.com/presentations/bh-usa-07/Ferguson/Whitepaper/bh-usa-07-ferguson-WP.pdf>

use after free

- using memory which is freed
 - leading to arbitrary code execution
 - common in c++
-
- <https://sploitfun.wordpress.com/2015/06/16/use-after-free/>

off by one

- **chunk overlapping**
 - extending free chunks
 - extending allocated chunks
 - shrinking free chunks
- unlink
- Glibc Adventures: The Forgotten Chunks
- https://www.contextis.com/documents/120/Glibc_Adventures-The_Forgotten_Chunks.pdf

- unsorted bin attack
- fast bin attack
-